


# HPy

## The Future of Python Native Extensions

**Florian Angerer**  
**Štěpán Šindelář**

Oracle Labs  
GraalPy and HPy Core Developers





# About Us and This Talk


- Who are we
  - Florian Angerer and Štěpán Šindelář
  - Software Engineers @ Oracle Labs
  - GraalPy core developers <https://graalvm.org/python>
  - HPy core developers <https://hpyproject.org>
    - not in group of HPy founders but joined very early in 2019
- This Talk
  - Motivates and introduces HPy
  - Shows benefits to you
  - Convince you to use HPy 😊





# Target Audience

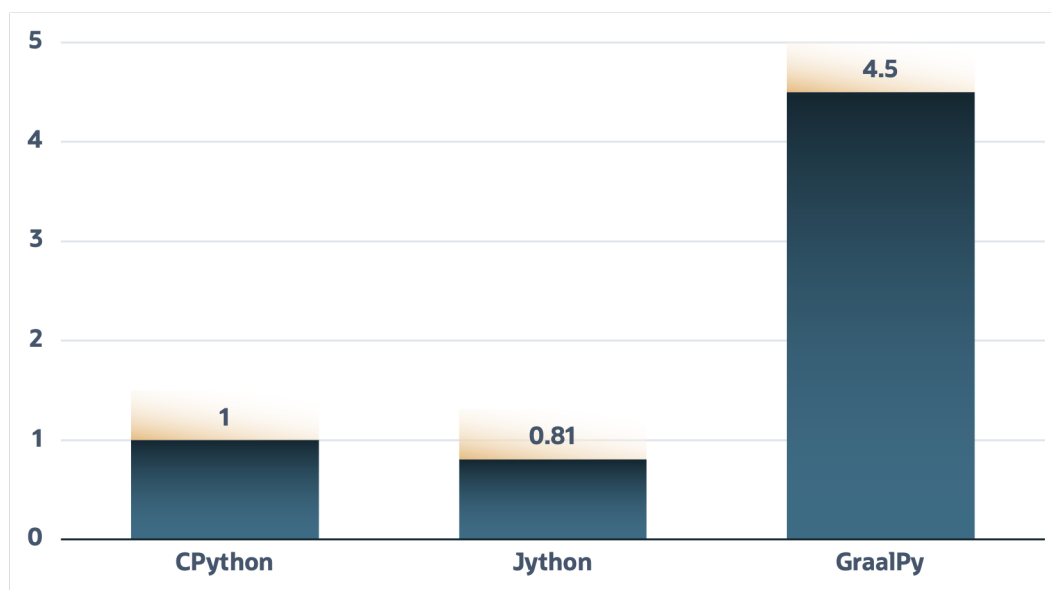
- “Experienced” Python programmer
  - Maybe even Python package developer
- You know C or the C memory model (pointers, struct, malloc, ...)
- Have a vague understanding of the (C)Python internals
- Have a vague understanding of `#include <Python.h>`



# CPython and Alternative Runtimes

- CPython
  - The Python reference implementation
  - A bytecode interpreter
  - Written in C
- Alternative Runtimes
  - GraalPy, PyPy, Jython, IronPython, Pyston, Cinder, ...
- Most of them try to improve Python language execution speed by using a just-in-time compiler, different data structures, moving GC, etc.

On average, GraalPy is 4.5x faster than CPython.




Geomean speedup over CPython on the [Python Performance](#) suite  
(Note that Jython can only run a subset of the benchmarks due to the missing Python 3 support)





# Python C API

- Since CPython is written in C
- Allow C extensions
- Exposes (internal) APIs, data structures, and implementation details
  - `ob_refcnt`, `ob_type`, `PyTypeObject`, ...
- Reference counting
- Objects as C pointers (`PyObject *`)
  - Exposes memory location
  - Assumption about identity
- Borrowed references
- Victor Stinner's PEP 620 <https://peps.python.org/pep-0620/>



## Example: Reference Counting vs. GC

- GraalPy is written in Java
- Java has the most advanced and mature GC implementations
- Reference counting prevents using a “real” GC
- The GC is NOT (only) about collecting garbage!
  - It should be called Memory Manager
- State-of-the-art GCs
  - super-fast allocation, fast deallocation
  - no/minimal pauses
  - multi-threading
- Why should I care about that?

# ~79x faster

GraalPy compared to CPython on gcbench.py

PyPy is also ~25x faster than CPython



# What is HPy?

- A novel C API for writing Python extensions
  - `#include "hpy.h"` instead of `"Python.h"`
- Funded directly/indirectly via OpenCollective, Oracle, IBM, Quansight Labs, and Anaconda Inc.
- Mostly driven by PyPy and GraalPy developers
- HPy: more abstract
  - Hides implementation details
  - Easier/faster on alternative implementations
  - GC friendly

ORACLE

 ANACONDA®

 open collective

IBM


 Quansight.  
YOUR DATA EXPERTS





# HPy Goals

- Zero overhead on CPython
  - Using macros and static inline functions to map HPy calls to C API calls
- Incremental migration path
  - Port existing extensions one function at a time
- Faster on alternative implementation
  - PyPy, GraalPy, ...
- Better debugging experience
- Universal ABI
  - One binary for multiple Python interpreters
- Backwards Compatibility
  - Multiple ABIs in the same interpreter



# How Does HPy Look?

```
// quickstart.c

#include "hpy.h" // instead of "Python.h"

HPyDef_METH(say_hello, "say_hello", HPyFunc_NOARGS)
static HPy say_hello_impl(HPyContext *ctx, HPy self)
{
    return HPyUnicode_FromString(ctx, "Hello world");
}

static HPyDef *QuickstartMethods[] = {
    &say_hello,
    NULL,
};

static HPyModuleDef quickstart_def = {
    .doc = "HPy Quickstart Example",
    .defines = QuickstartMethods,
};

HPy_MODINIT(quickstart, quickstart_def)
```

```
# setup.py
#
# python3 setup.py [--hpy-abi=universal] build

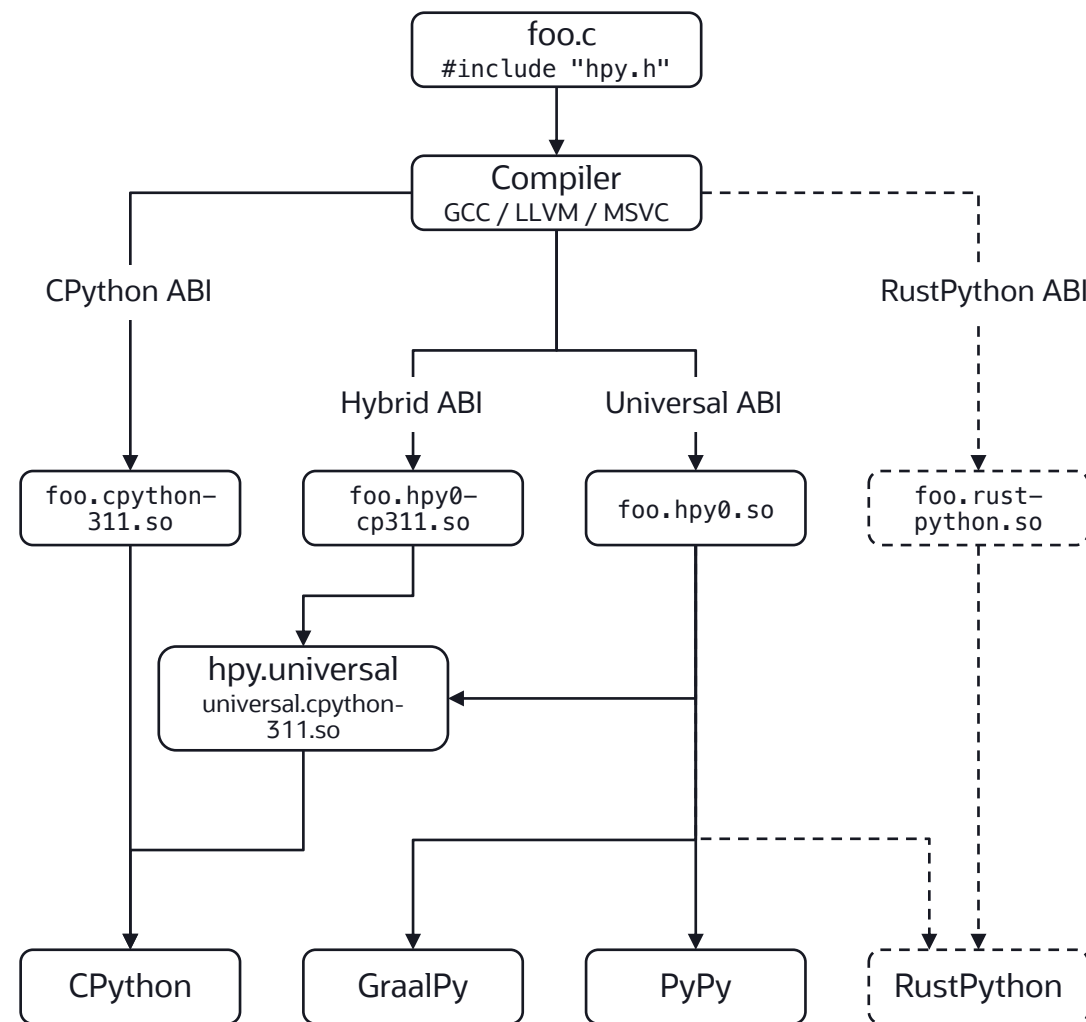
from setuptools import setup, Extension
from os import path

DIR = path.dirname(__file__)
setup(
    name='hpy-quickstart',
    hpy_ext_modules=[
        Extension('quickstart',
            sources=[path.join(DIR, 'quickstart.c')]),
    ],
    setup_requires=['hpy'],
)
```



# Zero Overhead on CPython

- Multiple compilation modes
- Universal ABI
  - Common ABI supported by all interpreters
  - One binary for all interpreters
  - API calls are done via HPyContext
- Custom ABIs
  - Maps HPy API functions to interpreter-specific APIs
  - E.g. CPython ABI
  - Binary is interpreter-specific
  - Best performance for target interpreter
- Hybrid ABI





# Incremental Migration Path

- Convert to HPy module
  - Keep existing functions as legacy methods
- Migrate types to (legacy) HPy type
- Migrate functions, slots, and members one by one
- Use `HPy_AsPyObject` and `HPy_FromPyObject`
- Build and test after each step with `--hpy-abi=hybrid`

```
static PyMethodDef PointModuleLegacyMethods[] = {
    {"dot", (PyCFunction)dot, METH_VARARGS, "Dot product."},
    {NULL, NULL, 0, NULL}
};

static HPyDef *module_defines[] = {
    &module_exec,
    NULL
};

static HPyModuleDef moduledef = {
    .doc = "Point module",
    .legacy_methods = PointModuleLegacyMethods,
    .defines = module_defines,
};
```


<https://docs.hpyproject.org/en/latest/porting-example/index.html>






# Performance

Kiwi solver port benchmark (8192 inner loops, 10 repetitions)



# Performance

## Numpy







# Debug Mode

- Another suitable name: Strict Mode
- Optional run-time mode (no recompilation required)
- Strictly enforces HPy contract
  - Does additional book keeping for resources
  - Prevents wrong API usage that happens to work by mistake on a certain implementation
- Examples
  - Leaked handles
  - Use-after-close
  - Lifetime of data pointers
  - Read-only memory
  - Invalid reuse of HPy context





# Debug Mode: Leak Detector

```
HPyDef_METH(test_leak, "test_leak", HPyFunc_NOARGS)
static HPy
test_leak_stacktrace_impl(HPyContext *ctx, HPy self)
{
    HPy num = HPyLong_FromLong(ctx, 42);
    if (HPy_IsNull(num)) {
        return HPy_NULL;
    }
    // missing HPy_Close(ctx, num);
    return HPy_Dup(ctx, ctx->h_None);
}
```

# Run with HPY=debug

```
import hpy.debug
import snippets

hpy.debug.set_handle_stack_trace_limit(16)
from hpy.debug import LeakDetector

with LeakDetector() as ld:
    snippets.test_leak_stacktrace()
```

```
Traceback (most recent call last):
  File "debug-example.py", line 7, in <module>
    snippets.test_leak_stacktrace()
  File "leakdetector.py", line 43, in __exit__
    self.stop()
  File "leakdetector.py", line 36, in stop
    raise HPyLeakError(leaks)

hpy.debug.leakdetector.HPyLeakError: 1 unclosed handle:
<DebugHandle 0x556bbcf907c0 for 42>
```

Allocation stacktrace:

```
universal.cpython-38d-x86_64-linux-
gnu.so(debug_ctx_Long_FromLong+0x45) [0x7f1d928c48c4]
snippets.hpy.so(+0x122c) [0x7f1d921a622c]
snippets.hpy.so(+0x14b1) [0x7f1d921a64b1]
universal.cpython-38d-x86_64-linux-
gnu.so(debug_ctx_CallRealFunctionFromTrampoline+0xca)
[0x7f1d928bde1e]
snippets.hpy.so(+0x129b) [0x7f1d921a629b]
snippets.hpy.so(+0x1472) [0x7f1d921a6472]
...
```

# Universal ABI



# DEMO





# NumPy/HPy

- A very hard one to migrate
- Why NumPy?
  - If we can do NumPy → we can do everything
- We invested ~1 year (FTE)
- ~180 kLOC (ANSI-C) and 80 kLOC (C API Usage)
  - ~40 kLOC changed
  - ~15 kLOC fully migrated
  - NumPy API: 118 / 261 (45 %)
- We think, hardest part has been done
  - Type migration
  - Metaclass support

<https://github.com/hpyproject/numpy-hpy/tree/graal-team/hpy>



# Current Status

- Current release: 0.9
- We've (\*partially) migrated several popular packages to HPy
  - ultrajson
  - matplotlib\*
  - psutils
  - kiwi solver
  - Pillow\*
  - Piconumpy
  - Numpy\*
- Cython backend is also planned (PoC exists)



# HPy Roadmap (Unordered List)

- Do stable release 0.9.0 (currently: 0.9.0rc2)
- Concentrate on NumPy/HPy
  - Migrate all types to heap types (no HPy involved yet)
  - Evaluate performance
  - Merge upstream
  - Continue migration to HPy
  - Evaluate performance
  - Merge upstream
- Release HPy 1.0.0





# Where We Need Help?

- Contributions in form of code, time, or money
- Documentation
  - Core devs concentrate on technical tasks
- Publicity
  - Migrate your package to HPy
  - Write a new package using HPy
- Tooling
  - Anything that helps using HPy
- Packaging
  - Integrate with build systems and PyPI
- Website and logo design



# Thank you

---



ORACLE