




EUROPYTHON — 17-23 July
PRAGUE & REMOTE 2023

Zero-Copy Zen

Boost Performance with
Memory View





EUROPYTHON PRAGUE & REMOTE

17-23 July
2023



Kesia Mary Joies
Product Engineer

U ·
S T



Aby M Joseph
Product Engineer

U ·
S T




Zero-copy



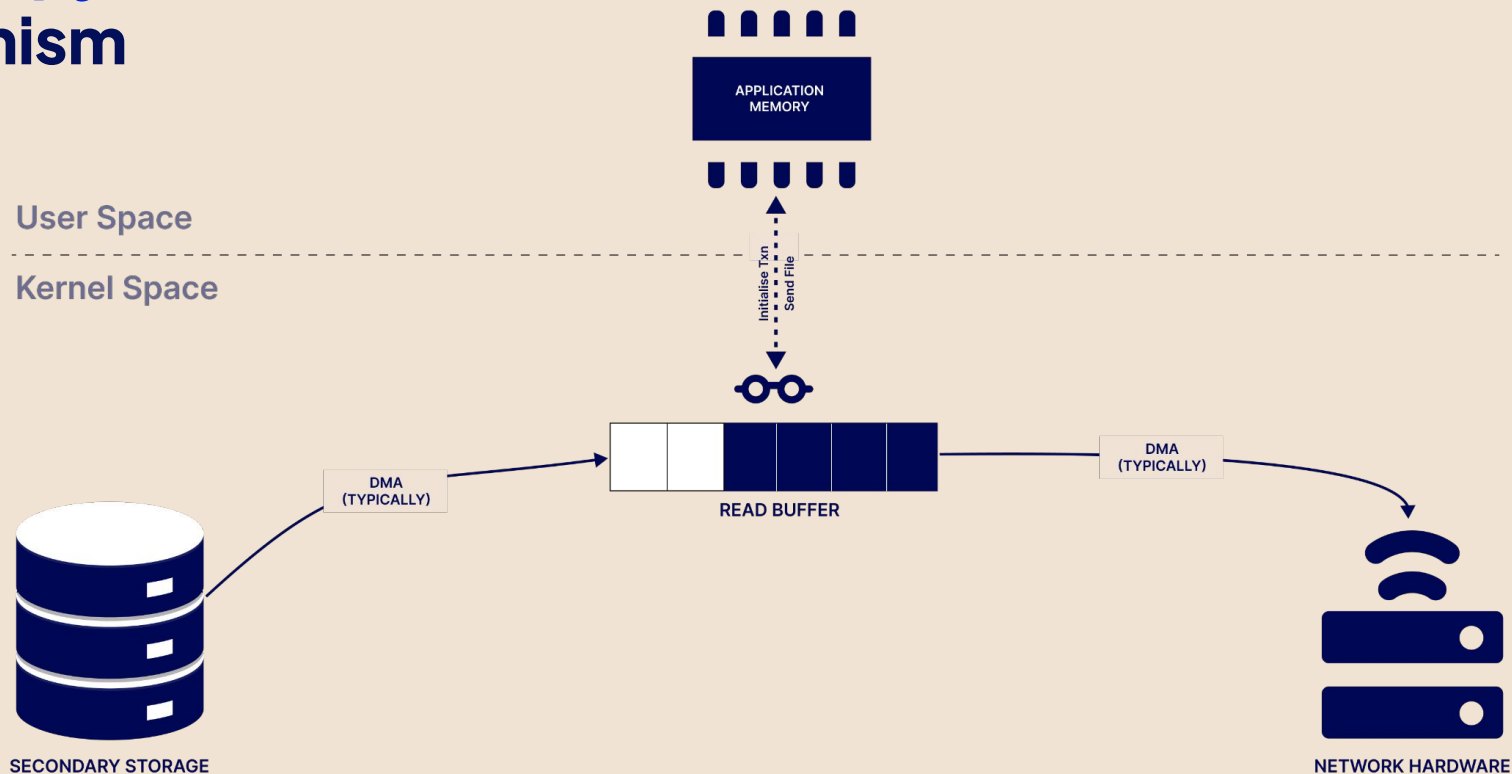
**A method to copy data from the
disk/network to the memory without
passing through the CPU**



Traditional Mechanism



Zero-copy Mechanism



Let's look into
an **example**



```
import socket
import hashlib

sock = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
server_port = 8082
server_addr = ("0.0.0.0", server_port)
print(f"Start server on port {server_port}")
sock.bind(server_addr)
sock.listen(1)

while True:
    print("Waiting for connection")
    connection, client_addr = sock.accept()
    size = 0
    try:
        i = 0
        while True:
            data = connection.recv(65536)
            i += 1
            if data:
                size += len(data)
            else:
                print("Done receiving data")
                break
        print(f"Total size: {size}")
    finally:
        connection.close()
```

Receiving server




```
import socket
import time

sock = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
server_port = 8082
server_address = ('127.0.0.1', server_port)
sock.connect(server_address)

start = time.time()
try:
    with open(r'/tmp/large_file', 'rb') as f:
        message = f.read()
        sock.sendall(message)
finally:
    sock.close()

end = time.time()
print('Total time: ', end-start)
```

Traditional Copy Client



```
import socket
import time

sock = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
server_port = 8082
server_address = ('127.0.0.1', server_port)
sock.connect(server_address)

start = time.time()
try:
    with open(r'/tmp/large_file', 'rb') as f:
        message = f.read()
        sock.sendall(message)
finally:
    sock.close()

end = time.time()
print('Total time: ', end-start)
```

```
"""
# Start server
$ python zerocopy_server.py
Start server on port 8082
Waiting for connection
Done receiving data
Total size: 2147483648
Waiting for connection
$ truncate -s 2G /tmp/large_file
$ python traditional_client.py
Total time: 2.4229979515075684
"""
```

A test run using a test file with **2GB** size





```
import os
import socket
import time

sock = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
server_port = 8082
server_address = ("127.0.0.1", server_port)
sock.connect(server_address)
start = time.time()

try:
    with open(r"/tmp/large_file", "rb") as f:
        ret = 0
        offset = 0

        while True:
            ret = os.sendfile(sock.fileno(), f.fileno(), offset, 65536)
            offset += ret
            if ret == 0:
                break

finally:
    sock.close()

end = time.time()
print("Total time: ", end - start)
```

Zero – Copy Client





```
import os
import socket
import time

sock = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
server_port = 8082
server_address = ("127.0.0.1", server_port)
sock.connect(server_address)
start = time.time()

try:
    with open(r"/tmp/large_file", "rb") as f:
        ret = 0
        offset = 0

        while True:
            ret = os.sendfile(sock.fileno(), f.fileno(), offset, 65536)
            offset += ret
            if ret == 0:
                break

finally:
    sock.close()

end = time.time()
print("Total time: ", end - start)
```



```
"""
# Server output
Waiting for connection
Done receiving data
Total size: 2147483648
Waiting for connection
...
# Zero copy client
$ python zerocopy_client.py
Total time: 1.161919116973877
"""
```



In Python, we can use Zero-copy
by using **memoryview**



Python's built-in types to manipulate binary data



Python's built-in types to manipulate binary data

Bytes



**Python's built-in
types to manipulate
binary data**

Byte Array

Bytes



Python's built-in types to manipulate binary data

Byte Array

Bytes

Memory View



Bytes

Immutable data type



Bytes

Return a byte object

Immutable data type



Bytes

Return a byte object

Immutable data type

Must be iterable of integers between $0 \leq x < 256$



Bytes

Syntax:

bytes([source[, encoding[, errors]])






bytes.py

```
data = bytes(4)
print(data)
# b'\x00\x00\x00\x00'


print(type(data))
# <class 'bytes'>
```





```
count = [1, 5, 4, 8, 2]
count_new = bytes(count)
print(count_new)
# b'\x01\x05\x04\x08\x02'
print(type(count_new))
# <class 'bytes'>
```






```
count = [1, 5, 4, 8, 2]
count_new = bytes(count)
print(count_new)
# b'\x01\x05\x04\x08\x02'
print(type(count_new))
# <class 'bytes'>

print(*count_new)
```






```
count = [1, 5, 4, 8, 2]
count_new = bytes(count)
print(count_new)
# b'\x01\x05\x04\x08\x02'
print(type(count_new))
# <class 'bytes'>

print(*count_new)
# 1 5 4 8 2
```






```
count = [1, 5, 4, 8, 2]
count_new = bytes(count)
print(count_new)
# b'\x01\x05\x04\x08\x02'
print(type(count_new))
# <class 'bytes'>

print(*count_new)
# 1 5 4 8 2

count_new[0] = 100
```





```
count = [1, 5, 4, 8, 2]
count_new = bytes(count)
print(count_new)
# b'\x01\x05\x04\x08\x02'
print(type(count_new))
# <class 'bytes'>

print(*count_new)
# 1 5 4 8 2

count_new[0] = 100
# TypeError: 'bytes' object does not support item assignment
```



Bytearray



Bytearray

Similar to bytes



Bytearray

Supports mutability

Similar to bytes





bytearray.py

```
data = bytearray(4)
print(data)
# bytearray(b'\x00\x00\x00\x00')

print(type(data))
# <class 'bytearray'>
```





bytearray.py

```
count = [10, 20, 15, 80, 53]
count_new = bytearray(count)
print(type(count_new))
# <class 'bytearray'>
```

```
print(*count_new)
# 10 20 15 80 53
```

```
count_new[0] = 100
print(*count_new)
# 100 20 15 80 53
```



Memory View

Zero-Copy View



Memory View

Slicing and Indexing

Zero-Copy View



Memory View

Slicing and Indexing

Zero-Copy View

Efficient Data Manipulation



Memory View

Syntax:
memoryview(object)





```
>>> view = memoryview(b'hello world')  
>>> view  
<memory at 0x0000025D2D26B1C8>
```





```
>>> view = memoryview(b'hello world')
>>> view
<memory at 0x0000025D2D26B1C8>

>>> view[0]
104
```





```
>>> view = memoryview(b'hello world')
```

```
>>> view
```

```
<memory at 0x0000025D2D26B1C8>
```

```
>>> view[0]
```

```
104
```

```
>>> chr(view[0])
```

```
'h'
```





memory_view.py

```
byte_array = bytearray('EURO', 'utf-8')  
mv = memoryview(byte_array)  
print(*mv)  
# 69 85 82 79
```



Converts to byte array
using str.encode()





memory_view.py

```
byte_array = bytearray('EURO', 'utf-8')
mv = memoryview(byte_array)
print(*mv)
# 69 85 82 79

print(type(mv))
# <class 'memoryview'>
```



Memory View

Uses buffer protocol



So, what is buffer protocol ?



Buffer Protocol

It is a protocol that provides a way to access the internal data of an object.



Whenever we perform some action on an object,
Python needs to create a **copy** of the object.





list_slice.py

```
my_list = [1, 2, 3, 4, 5]
print(f'{my_list=}')
# my_list=[1, 2, 3, 4, 5]

sliced_list = my_list[1:3]
print(f'{sliced_list=}')
# sliced_list=[2, 3]
```



Buffer Protocol

But the problem is, even if we *cannot access this protocol* with the standard code base, this is accessible to us at the C-API level.



Buffer Protocol

But the problem is, even if we *cannot access this protocol* with the standard code base, this is accessible to us at the C-API level.

So, in Python, if we want to expose the same protocol in Python, we *need to use memoryview*.



Buffer Protocol

Improves execution speed



Buffer Protocol

Use less memory

Improves execution speed



Buffer Protocol

Use less memory

Improves execution speed

Works on large data



PEP 688 – Making the buffer protocol accessible in Python

- A better way to interact with the buffer protocol directly.
- Now supports additional types with the buffer protocol.
- Simplified Workflow
- Performance Boost



Diving into Memory View





memory_view.py

```
s1 = b"Hello World"
s2 = bytearray(b'Hello World')

s1View = memoryview(s1)
s2View = memoryview(s2)

print("Bytes is readonly?:", s1View.readonly)
# Bytes is readonly?: True
print("Byte array is readonly?:", s2View.readonly)
# Byte array is readonly?: False
```





memory_view.py

```
# Modifying object through writeable view
```

```
s2 = bytearray(b'Hello World!')
```

```
s2View = memoryview(s2)
```

```
print("Before:", s2)
```

```
# Before: bytearray(b'Hello World!')
```





memory_view.py

```
# Modifying object through writeable view
```

```
s2 = bytearray(b'Hello World!')  
s2View = memoryview(s2)
```

```
print("Before:", s2)  
# Before: bytearray(b'Hello World!')
```

```
s2View[6:12] = b'Python'  
print("After:", s2)  
# After: bytearray(b'Hello Python')
```





```
import array
a = array.array('l', [-11111111, 22222222, -33333333, 44444444])
m = memoryview(a)
print(m[0])
# -11111111
print(m[-1])
# 44444444
print(m[::2].tolist())
# [-11111111, -33333333]
```



Comparison



```
import time

print("Type\t\t\tIterations\t\tTime Taken")
print("-----")

for n in (100000, 200000, 300000, 400000, 500000):
    data = b'x'*n
    start = time.time()
    b = data
    while b:
        b = b[1:]
    end = time.time()
    tm = end - start
    print(f'bytes \t\t\t{n}\t\t\t {tm:0.2f}')
```

Without Memory View



```
import time

print("Type\t\t\tIterations\t\tTime Taken")
print("-----")

for n in (100000, 200000, 300000, 400000, 500000):
    data = b'x'*n
    start = time.time()
    b = data
    while b:
        b = b[1:]
    end = time.time()
    tm = end - start
    print(f'bytes \t\t\t{n}\t\t\t\t {tm:0.2f}')
```

Without Memory View

```
import time


print("Type\t\t\t\tIterations\t\tTime Taken")
print("-----")

for n in (100000, 200000, 300000, 400000, 500000):
    data = b'x'*n
    start = time.time()
    b = memoryview(data)
    while b:
        b = b[1:]
    end = time.time()
    tm = end - start
    print(f'memoryview \t\t\t{n}\t\t\t\t {tm:0.2f}')
```

With Memory View



Without Memory View




```
"""
```

Type	Iterations	Time Taken
bytes	100000	0.24
bytes	200000	0.91
bytes	300000	1.62
bytes	400000	2.64
bytes	500000	4.33

```
"""
```

With Memory View

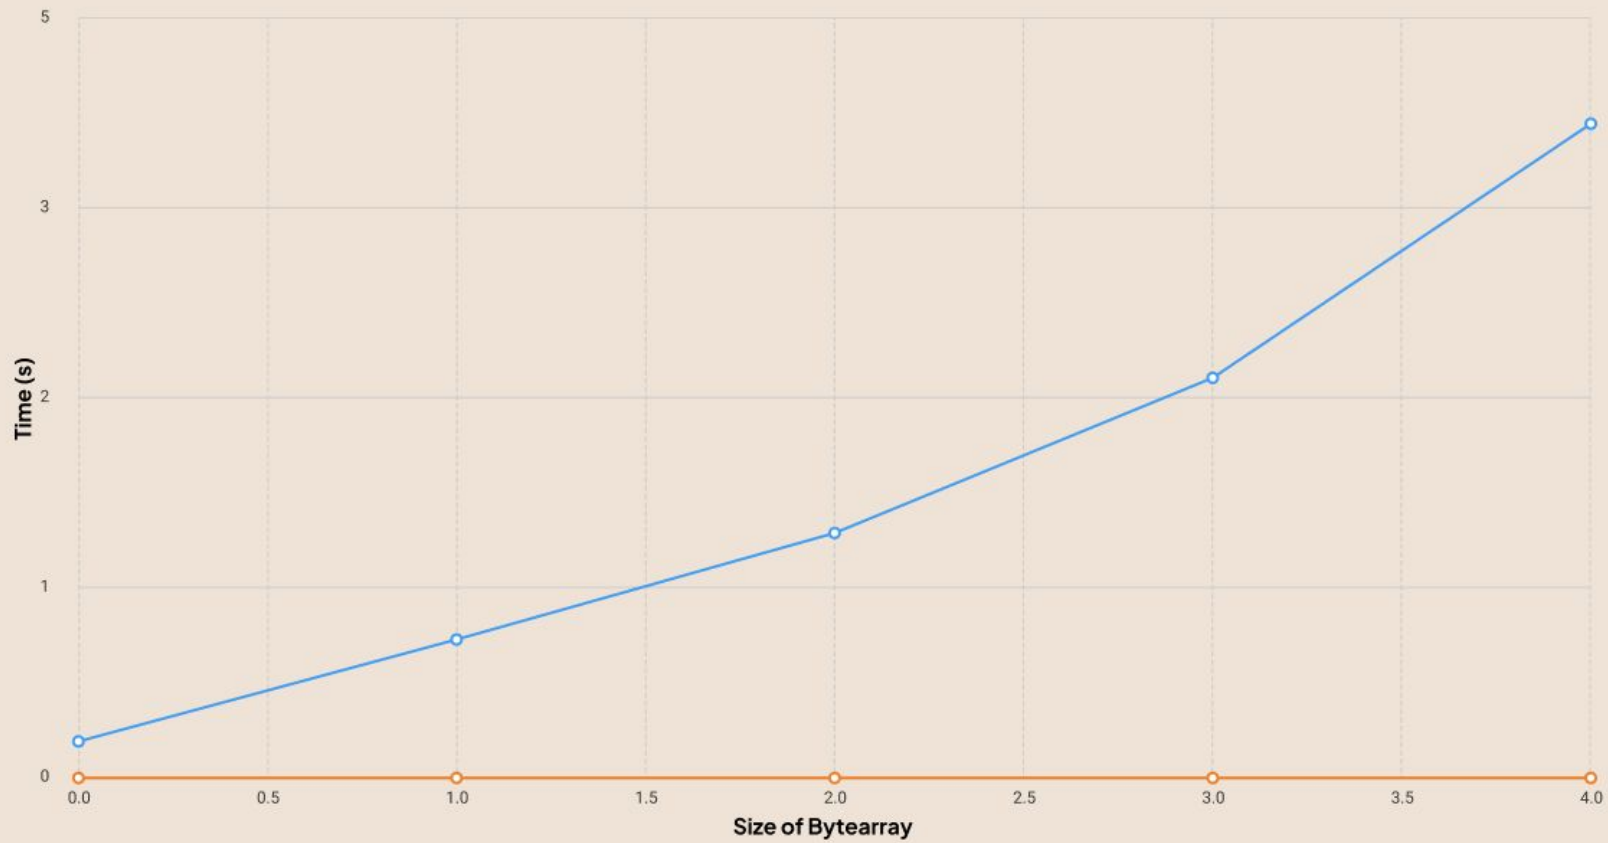


```
"""
```

Type	Iterations	Time Taken
memoryview	100000	0.02
memoryview	200000	0.03
memoryview	300000	0.05
memoryview	400000	0.07
memoryview	500000	0.08

```
"""
```





Applications



Example – Numeric and Scientific Computing:



```
import numpy as np

# Create a large NumPy array
size = 1000000
large_array = np.arange(size)

# Create a memory view of the array
mem_view = memoryview(large_array)

# Double the values in the array using memory view
for i in range(len(mem_view)):
    mem_view[i] *= 2

# Print the first 10 elements of the modified array
print(large_array[:10])
# [ 0  2  4  6  8 10 12 14 16 18]
```



Example – Video Streaming

```
def timecode_to_index(video_id, timecode):  
    ...  
    # Returns the byte offset in the video data  
  
def request_chunk(video_id, byte_offset, size):  
    ...  
    # Returns size bytes of video_id's data from the  
    offset  
video_id = ...  
timecode = '01:09:14:28'  
byte_offset = timecode_to_index(video_id, timecode)  
size = 20 * 1024 * 1024  
video_data = request_chunk(video_id, byte_offset, size)
```



Example – Video Streaming



```
import timeit

def run_test():
    chunk = video_data[byte_offset:byte_offset + size]
    # Call socket.send(chunk), but ignoring for
    benchmark
    result = timeit.timeit(
        stmt='run_test()',
        globals=globals(),
        number=100) / 100

print(f'{result:0.9f} seconds')
>>>
0.004925669 seconds
```



Example – Video Streaming



```
video_view = memoryview(video_data)

def run_test():
    chunk = video_view[byte_offset:byte_offset + size]
    # Call socket.send(chunk), but ignoring for
    benchmark
    result = timeit.timeit(
        stmt='run_test()',
        globals=globals(),
        number=100) / 100

print(f'{result:0.9f} seconds')
>>>
0.0000000250 seconds
```



Benefits

Reduced Memory Footprint



Benefits

Better Performance

Reduced Memory Footprint



Benefits

Better Performance

Reduced Memory Footprint

Improved Scalability



References

1. [Idea of zero copy \[with example\]](#)
2. [Effective Python › Item 74: Consider memoryview and bytearray for Zero-Copy Interactions with bytes](#)
3. [Python memoryview\(\)](#)
4. [Python Memoryview Example for Beginners](#)
5. [memoryview in Python](#)
6. [Python memoryview\(\) Function \(Buffer Protocol And Memory View\)](#)
7. [PEP 688 – Making the buffer protocol accessible in Python](#)





EUROPYTHON — 17-23 July
PRAGUE & REMOTE 2023

Thank you

@kesiajoies 

@abymjoseph 