

Diving into Event-Driven Architectures with Python

Solving complexity at scale

EuroPython 2023, Prague, Czech Republic – 21.07.2023

Marc-André Lemburg :: eGenix.com GmbH

Speaker Introduction


- Marc-André Lemburg
 - Python since 1994
 - Studied Mathematics
 - CEO **eGenix.com GmbH**
 - **Senior Solution Architect, Consulting CTO and Coach**
 - **Python Core Developer (PEP 100, DB-API)**
 - **EuroPython Society Fellow** and former chair
 - **Python Software Foundation Fellow** and former director
 - **Co-founder Python Meeting Düsseldorf**
 - Based in Düsseldorf, Germany
 - More details: <http://malemburg.com>



Motivation: You're thinking of building the next Big Thing ...




Motivation: ... but you have no idea where it will take you



Motivation: Prepare for growth and flexibility

- Your architecture needs to do be:
 - **scalable (horizontally and vertically)**
 - easily adapt to new challenges
 - easy to maintain for devops
 - **prepared for the enterprise**
- It should also:
 - have good failure modes
 - **integrate observability**
 - automate governance



... just to name a few goals 😊

Well-known synchronous architectures

- **REST**

- Good frontend support
- Many available backend systems
- Complex when it comes to adapting to new data
- Difficult to scale



- **GraphQL**

- Simplifies querying new data
- New backend systems being developed, frontend support progressing
- Difficult to scale

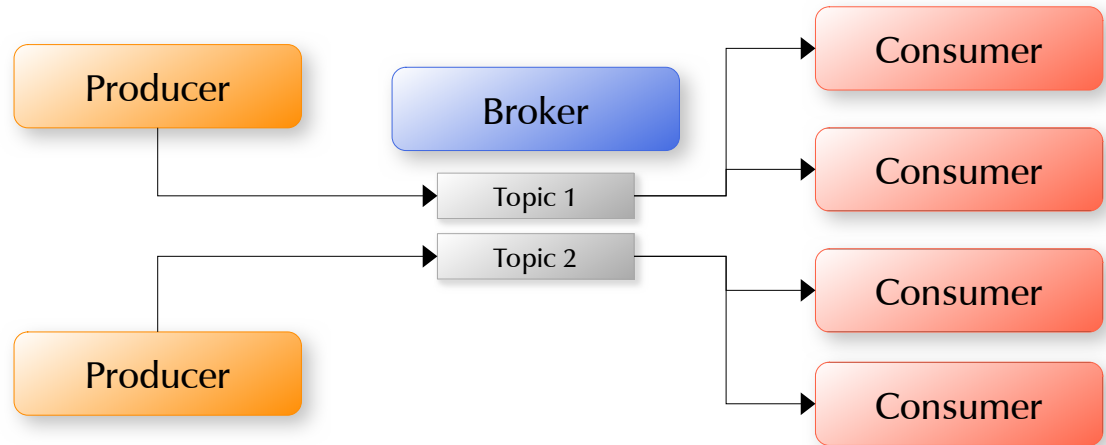
Event-Driven Architectures – *Rediscovered*

- EDA
 - Asynchronous
 - Can be combined with REST and GraphQL
 - Data agnostic
 - Scales well and promotes loose coupling
 - (Potentially) Solves many of the problems with integrating complex systems
 - Originally from the early 2000s
- *Let's have a closer look...*



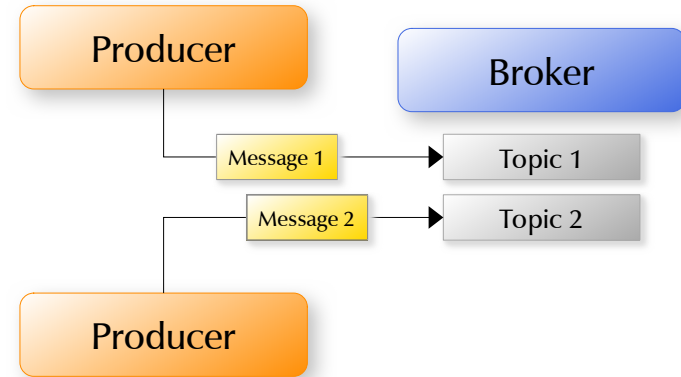
Main concept: Event-driven Communication

- Events are produced and consumed **without direct connection**
- Events are organized in **topics**
- A **broker** manages event distribution
- Fully **asynchronous**
- *Note: Consumers can be Producers as well*



Event Messages

- Events are captured using **messages**
 - “something happened”
 - “state changed”
 - “initiate command”
- Messages
 - Headers for meta data
 - **Encoded payload**
 - Best practice: 100-1000 bytes per message
- **Bulk data** is better stored in an object store or database



Event Message Payloads

- Common formats:
 - **Apache Avro** (default for Kafka)
 - Protobuf
 - Apache Thrift
 - MessagePack
 - JSON (typed using JSON Schema)
- **Should be typed and signed** for better security
- **Compression** helps keep traffic reasonable
 - zlib
 - snappy

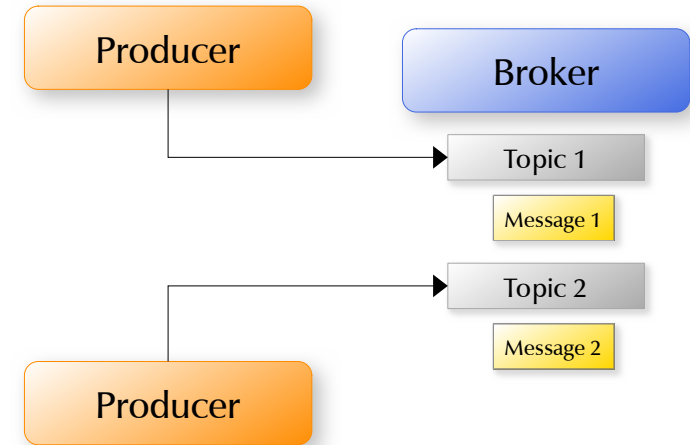


MessagePack



Event Message Distribution (1/2)

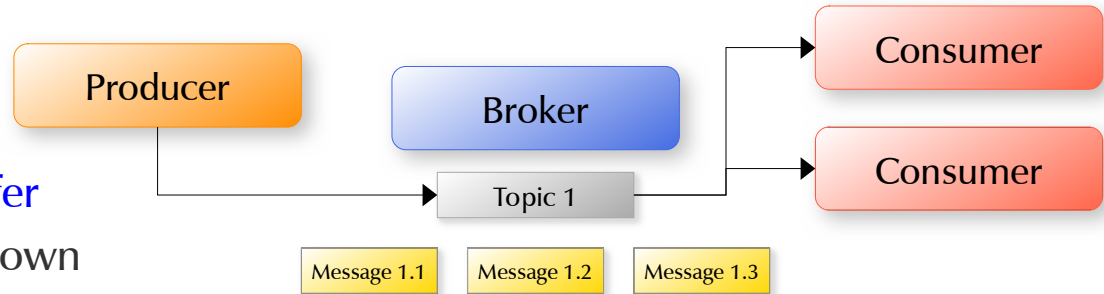
- **Topics**
 - Used to organize messages
 - Best practice: **one message type per topic**
 - Other terms: *channels, queues*
- **Publish/Subscribe (PubSub)**
 - Producers publish messages to a topic, consumers subscribe to messages on a topic
 - Brokers only **queue messages until delivered**
 - **Messages cannot be resent to new subscribers**



Event Message Distribution (2/2)

- **Streaming**

- Broker stores messages in a **stream buffer**
- Consumers manage their own reading from the stream
- **Enables replay and late joining**




- **Decoupled sending / receiving**

- Producer doesn't know who will be receiving the messages
- Consumer doesn't need to know who produced the messages
- **Separation of concerns**
- **Encapsulation**

Message Brokers

- Custom protocols
 - **Apache Kafka**, Apache Pulsar, PostgreSQL, Redis
- Focus: **AMQP**
 - **RabbitMQ**
- Focus: **MQTT**
 - Apache ActiveMQ, HiveMQ, Mosquitto
- Cloud
 - AWS SNS, Google PubSub, Azure PubSub

- Older variants
 - **IBM MQSeries**




- Often come with **connectors** to simplify integrating sources (producers) and sinks (consumers)
 - Kafka Connect
 - Pulsar Connectors



Broker Challenges

- **Guaranteed delivery of messages**
 - Even in the face of network issues, failures, etc.
- **Processing messages exactly once**
 - Make processing idempotent to soften this requirement
- **Failure modes**
 - Automatic retry/replay in case of failures
 - Backfilling data
- **Trade-off message size vs. performance**
 - Small message often result in additional source data queries
 - Larger messages are slower to handle and store



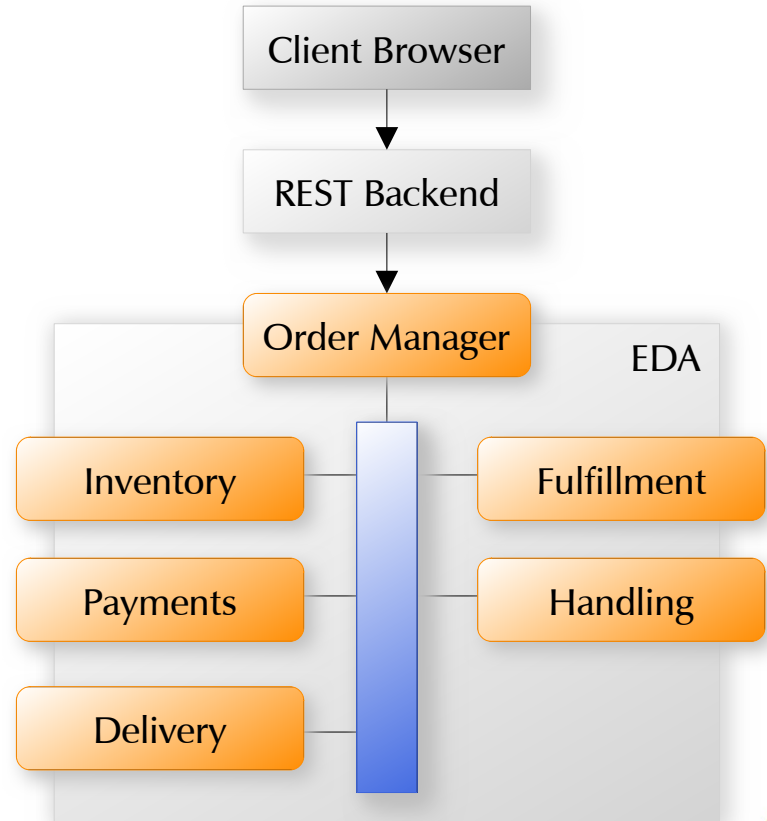
EDA API Specification and Documentation

- Requirements
 - Needed in machine readable form
 - Needed for discovery and ease of use
 - Needed for secure event API interaction (type checking)
- AsyncAPI – <https://asynccapi.com/>
 - Machine readable
 - Similar to OpenAPI/Swagger (REST)
 - Adapted for async processing, pub/sub
 - Relatively new (started in 2017)
 - Python is not really a key player (yet)
 - Node.js and Java dominate



Event-Driven Architectures (EDA) – A Recap


- Split applications into loosely coupled components
 - Have components communicate via events
 - Operate in asynchronous mode
- Use a broker to manage communication
- Scale up/down individual parts of the stack as needed
 - even on-demand...



Comparison: REST/GraphQL vs. EDA

- *Let's have a look at an example*

Handle an order in web shop



Example: Handle order in web shop

- REST/GraphQL (Sync):
 - Interface to lots of subsystems to initiate order
 - Payment system
 - Fulfillment system
 - Inventory system
 - Queue order in handling and delivery system
 - Challenges
 - Shop backend is single point of failure (for the order)
 - Keeping track of dependencies between subsystems
 - Changes to a subsystems will often require changes to shop backend
 - Handling interfacing issues: network problems, temporary failures
 - Rollback order in case of problems



Example: Handle order in web shop

- EDA (Async):
 - Shop backend sends order event
 - Subsystems can react by processing their respective parts
 - An **order management system** manages the order
 - Assures successful completion of all subevents
 - Rolls back order in case of problems
 - **System load can easily be distributed and scaled**
 - Subsystem scaling can be applied independently
 - **Code management is distributed as well**
 - Fewer dependency issues
 - Easier to handle upgrades without changes to shop system



EDA sounds too good to be true: Challenges

- **Debugging Challenges**
 - Gathering logs from all nodes
 - Associating log entries with incoming requests
 - Introspection for processed requests
- **Organizational Challenges**
 - Common understanding of the architecture
 - Communicate the design to all participants
 - Document events, event types
 - AsyncAPI
 - Document expectations
 - When to send events
 - What to expect as a result



EDA and Python

- Python async support is great for pub/sub style APIs
 - Scales well
 - Easy to use
 - Can be combined with multi-threading access to interface libraries
 - External C libraries don't need the Python GIL
- *Let's have a closer look:*
 - *AsyncAPI support*
 - *Roll-your-own*
 - *Conclusion*



Python AsyncAPI Support

- Not much available for Python :-(
 - Would be great to get some more attention and support from the Python community
- AsyncAPI community
 - Much focus on Node.js and Java
- AsyncAPI Studio
 - Online tool for working with AsyncAPI specs
 - Can generate Paho (MQTT) code from spec
 - Just basic support for Python
 - Code generation not optimal



Python AsyncAPI Support: PyPI packages

- asyncapi package
 - Dynamically reads AsyncAPI spec and provides pub/sub APIs
 - Uses broadcaster package for the lower level pub/sub interfaces
 - Supports Redis, Kafka, PostgreSQL, Google pub/sub
 - Development stalled (last active in 2020)
- fastapi-asyncapi package
 - Almost no documentation
 - Helps expose AsyncAPI specs on a web service
 - Development stalled (last active in 2021)



Python Roll-your-own EDA: Low level PyPI packages

- Use available package to interface to brokers directly
 - **Kafka**: aiokafka
 - Pulsar: pulsar
 - Redis: asyncio-redis
 - Postgres: asyncpg
 - MQTT:
 - paho-mqtt (sync)
 - asyncio-paho (async)
 - RabbitMQ:
 - amqp or rabbitmq-client (sync)
 - aioamqp (async)
 - *(many more, e.g. for cloud services, etc.)*



EDA and Python: Conclusion

- Roll-your-own is definitely possible now
 - And used a lot in major companies
- We need better support for EDA in Python
 - More abstractions
 - Better integration into existing web backends
- More integration of AsyncAPI standards would also help
 - Dynamic spec generation
 - Dynamic API generation
 - Code generation tools



Main takeaway: Never stop to **learn** and **try out new things...**



These
are
exciting
times

Thank you for your attention !



Time for discussion

Contact

eGenix.com Software, Skills and Services GmbH

Marc-André Lemburg

Pastor-Löh-Str. 48

D-40764 Langenfeld

Germany

Email: mal@egenix.com

Phone: +49 211 9304112

Fax: +49 211 3005250

Web: <https://www.egenix.com/>

LinkedIn:



References

- Several photos taken from Pixabay and Unsplash
- Some screenshots taken from the mentioned websites
- All other graphics and photos are (c) eGenix.com or used with permission
- Details are available on request
- Logos are trademarks of their respective trademark holders